

THE SUMMIT DEBATES: Chiplet Summit 2026

Santa Clara Convention Center

A newcomer-friendly, first-principles debate on:

- Chiplet partitioning strategy for inference workloads
- A strategy for chiplet partitioning” (IO chiplet vs keeping IO on logic/memory, avoiding advanced packaging, etc.).

Panelists will defend real design choices: latency vs throughput targets, power/thermal envelopes (datacenter vs edge), bandwidth density vs yield (2.5D interposer vs 3D stacked memory), and IO placement.

The session ends with a one-page decision checklist attendees can take back to their teams.

Discussion (shared in advance)

- For a strict p99 latency target at edge power envelopes, when do chiplets hurt vs help?
- If memory bandwidth is the bottleneck, what tips the balance between 2.5D and 3D?
- IO: separate die vs co-locate with compute — what did you measure, and what broke?
- Validation surface and bring-up: how did chiplets change your schedule/risk profile?
- TCO: where did yields, packaging cost, and re-use of die hit the breakeven?

Target panel composition

- Two “chiplet-forward” implementers (startup and established).
- One “monolithic-leaning” voice (or neutral system perspective).
- One packaging/foundry expert (practical limits, yield, and cost).
- Moderator focused on constraints and decision criteria.



Volunteers	Email
Frank Schirrmeister	frank@schirrmeister.com
Joshua Rubin	jmrubin@us.ibm.com
Nick Ilyadis	#">nick@ilyadis.com #
Brian Reh	brian.h.rea@intel.com
Jim Finnegan	jim.finnegan@netronome.com
Paul Borrill	paul@daedaelus.com

Organizers	Email
Lance Leventhal	lanclev@pacbell.net
Elizabeth Leventhal	elizabeth@chipletsummit.com
Chuck Sobey	chuck@daedaelus.com

Deliverables and success metrics

- Handout: “Monolithic vs Chiplet — a decision checklist for inference.”
- Slide posted publicly: “Five claims tested; three that held up; two that didn’t.”
- Attendance target: top-quartile session for the track (prime slot + pre-event promotion).
- Post-session: publish the checklist and a short write-up to drive ongoing traffic to Chiplet Summit channels.

Core focus

- Partitioning strategies for inference workloads across chiplets.
- Latency vs throughput design trade-offs (and balanced designs).
- Deployment context differences: data center vs edge.
- 3D chiplet architectures (memory-on-compute and compute-on-memory) emerging fast.

Target speakers/companies

- Meta — recent publications on 3D memory for AR & LLMs.
- AMD — MI300/MI350 used for inference.
- Fujitsu — 3D chiplet-based AI product with Broadcom.
- Broadcom — 3D XDSiP platform backing multiple big players.
- Google — TPU v7 (Ironwood).
- Startups — d-Matrix, Tenstorrent, Persimmon(s) for breakout perspectives.
- Rationale
- AI hardware is hot; audience draw should be strong.
- Bias toward action: better to try and iterate than hesitate.

[Josh] Open to brainstorming another related session if we want to expand the track. When you’re ready for action items (outreach grid, draft invites, or panel structure), say the word and I’ll spin them up.

Candidate companies (initial list):

- AMD — attended Chiplet Summit last year (facilitated by Joshua).
- d-Matrix — shipping chiplet-based accelerator.
- Tenstorrent — chiplet roadmap.
- FuriosaAI — chiplet roadmap.
- Intel — past (Ponte Vecchio) + future plans; currently “not in a good position.”
- Meta — research with 3D prototypes (not explicitly branded “chiplet”).
- Observation: Today’s pool is thin, but Joshua expects rapid expansion within a year given rumors of hyperscaler silicon (Amazon, Google, Meta, Microsoft, OpenAI, Apple) with Broadcom/
- Marvell and 3.5D XDSiP / 3D SoIC packaging directions.
- Note: Supporting IP/providers (ARM, Alphawave, etc.) intentionally excluded here since they’ll be well represented elsewhere at the Summit.

This panel will bring together CTOs from across the compute, memory, and networking domains to explore:

Track organizers (current leads):

- AI accelerators: Joshua Rubin, Paul Borrill, Nick Ilyadis
- Die-to-die interfaces: Brian Rea, Mayank Bhatnagar
- Edge for AI: Feng Zhou, Dave Eggleston (interest: Jim Finnegan)
- HBM: Shahab Ardalan (also listed under Optical/Memory)
- Working with Foundries (Integration): Milind Weling, Nick Ilyadis
- Packaging: Trent Uehling, Nick Ilyadis
- Interest areas: Optical (Shahab), Test (Adam Cron); Soheil Modirzadeh can nominate for Packaging/Integration; Marc Hutner topic TBD.