

Pipelined Model Parallelism and Graph Languages for Infrastructure Management

Research Report

June 28, 2025

This research explores pipelined model parallelism (PMP) and examines existing graph representation standards to inform the design of a new graph language for infrastructure management. The study reveals significant opportunities for standardization in computational graph representations, particularly for distributed systems that employ pipeline parallelism.

Executive Summary

Pipeline model parallelism has emerged as a critical technique for training large-scale machine learning models, with sophisticated scheduling algorithms and optimization strategies. Current graph languages focus primarily on static representations rather than dynamic computational flows, creating a significant gap in infrastructure management capabilities.

The research identifies three primary gaps: (1) existing graph languages (DOT, GraphML, GQL) lack temporal and resource modeling capabilities, (2) no unified standard exists for representing pipelined computational infrastructures, and (3) current tools address specific domains without comprehensive infrastructure abstraction. These findings are detailed in Section 3 and inform the design recommendations presented in Section 5.

Key Findings:

- No unified standard for pipelined computational infrastructures
- Existing tools lack comprehensive infrastructure abstraction
- Significant optimization opportunities exist

Understanding Pipelined Model Parallelism

Foundational Concepts

Pipeline parallelism improves both the memory and compute efficiency of deep learning training by partitioning the layers of a model into stages that can be processed in parallel [1]. This approach addresses the fundamental challenge of training increasingly large models that exceed single-device memory capacity.

The technique employs three core principles:

Model Partitioning Splits neural network layers across multiple devices, with each device responsible for a subset of layers.

Gradient Accumulation Each backward pass accumulates gradients locally, followed by parallel gradient reductions across data parallel groups.

Microsoft DeepSpeed Team. *Pipeline Parallelism - DeepSpeed*. 2024. <https://www.deepspeed.ai/tutorials/pipeline/>

The core innovation of pipeline parallelism is the division of models into stages, enabling parallel processing while managing memory constraints effectively. See DeepSpeed documentation at <https://www.deepspeed.ai/tutorials/pipeline/>.

Micro-batch Processing Large training batches are divided into smaller micro-batches that flow through the pipeline stages.

Pipeline Scheduling Strategies

The research reveals two primary scheduling approaches with distinct trade-offs between efficiency and consistency [2, 3].

Synchronous Scheduling (GPipe) maintains gradient consistency but suffers from pipeline bubbles [4]. The focus is on balancing arithmetic intensity through micro-batch sizing while minimizing pipeline bubbles.

Asynchronous Scheduling (PipeDream) reduces pipeline stalls but introduces gradient staleness [2]. This approach revisits model parallelism for performance rather than memory limitations.

Advanced Optimizations

Modern pipeline implementations incorporate sophisticated optimizations:

Memory Management involves careful orchestration of activation memory communication between pipeline stages [1]. Each stage completes forward passes before communicating activations to subsequent stages.

Communication Optimization leverages the fact that pipeline-parallel training communicates significantly less data than data-parallel training, requiring only activation and gradient transfers at stage boundaries [5].

Load Balancing addresses the NP-hard problem of mapping computational graphs to resources, utilizing heuristic approaches for practical solutions [6].

Current Graph Language Landscape

Existing Standards Analysis

The analysis of current graph representation standards reveals significant limitations for infrastructure modeling.

DOT Language serves as the foundation for Graphviz visualization [7, 8]. While human-readable and widely supported, it lacks temporal representation, metadata support, and resource modeling capabilities. *DOT was historically an

Narayanan et al. *PipeDream: Generalized Pipeline Parallelism for DNN Training*. SOSP 2019.
Jang. *Parallelism in Distributed Deep Learning*. 2022. <https://insujang.github.io/2022-06-11/parallelism-in-distributed-deep-learning/>

Pipeline bubbles represent idle time when computational resources are underutilized during pipeline execution transitions.

Microsoft Research. *PipeDream: A more effective way to train deep neural networks using pipeline parallelism*. 2019. <https://www.microsoft.com/en-us/research/blog/pipedream-a-more-effective-way-to-train-deep-neural-networks-using-pipeline-parallelism/>
Computer Science Stack Exchange. *Scheduling distributed computational graph*. 2015. <https://cs.stackexchange.com/questions/50243/scheduling-distributed-computational-graph>

Wikipedia. *DOT (graph description language)*. 2024. [https://en.wikipedia.org/wiki/DOT_\(graph_description_language\)](https://en.wikipedia.org/wiki/DOT_(graph_description_language))
Graphviz Team. *DOT Language*. 2024. <https://graphviz.org/doc/info/lang.html>

acronym for "DAG of tomorrow," succeeding earlier DAG formats with broader graph structure support.

GraphML provides an XML-based format developed through joint efforts of the graph drawing community [9]. Despite extensible attributes and complex graph type support, it suffers from XML overhead and lacks built-in infrastructure concepts.

GQL (Graph Query Language) represents the first new ISO database standard since SQL in 1987 [10, 11]. While offering property graph focus and query capabilities, it remains database-oriented rather than infrastructure-specific.

Infrastructure-Specific Requirements

Current graph languages lack essential features for infrastructure representation across three critical dimensions.

Temporal Dynamics are absent from existing standards, yet pipeline stages require execution order, timing constraints, and resource availability modeling over time.

Resource Modeling capabilities are insufficient for representing memory, compute, and network bandwidth constraints across heterogeneous device capabilities [12].

Execution Semantics lack support for dataflow dependencies, error handling mechanisms, and performance optimization hints essential for modern infrastructure.

Infrastructure Management Context

Current Infrastructure Tools

Modern infrastructure management relies on diverse tools with varying graph representation approaches.

Kubernetes Architecture employs cluster-based models with control planes and worker nodes [13]. Visualization of Kubernetes deployments as graphs enables better understanding of component interactions [14].

Workflow Orchestration platforms like Apache Airflow utilize directed acyclic graphs (DAGs) for task scheduling and dependency management [15, 16].

Wikipedia. *GraphML*. 2024. <https://en.wikipedia.org/wiki/GraphML>

ISO/IEC. *ISO/IEC 39075:2024 - Database languages - GQL*. 2024. <https://www.iso.org/standard/76120.html>

Jackson. *GQL: A New ISO Standard for Querying Graph Databases*. The New Stack, 2024. <https://thenewstack.io/gql-a-new-iso-standard-for-querying-graph-databases/>

Neo4j Inc. *Graph DB for network database infrastructure monitoring*. 2024. <https://neo4j.com/use-cases/network-and-it-operations/>

Kubernetes Contributors. *Cluster Architecture*. 2024. <https://kubernetes.io/docs/concepts/architecture/>

The New Stack. *Reframing Kubernetes Observability with a Graph*. 2023. <https://thenewstack.io/reframing-kubernetes-observability-with-a-graph/> Kubernetes increasingly adopts graph-based observability for dependency management and impact analysis.

Apache Airflow. *Dags - Airflow Documentation*. 2024. <https://airflow.apache.org/docs/apache-airflow/stable/core-concepts/dags.html>

Apache Airflow. *What is Airflow?* 2024. <https://airflow.apache.org/docs/apache-airflow/stable/index.html>

Infrastructure as Code movement toward standard programming languages enhances modularity and maintainability [17].

Computational Graph Frameworks

Existing ML frameworks provide insights into effective graph representation patterns.

TensorFlow employs dataflow graphs representing computation relationships between operations [18, 19]. These graphs enable parallel execution, optimization, and portability across environments.

ONNX establishes an open standard for ML model interoperability [20, 21]. The framework defines common operators and file formats for cross-framework compatibility.

Design Recommendations

Core Language Requirements

Based on research findings, a comprehensive infrastructure graph language should incorporate multi-level abstraction capabilities:

Physical Layer Hardware topology and network connections

Logical Layer Services, containers, and virtual resources

Application Layer Workloads, data flows, and dependencies

Temporal Layer Schedules, lifecycles, and state transitions

Pipeline-Aware Constructs

The language must provide native support for:

- Stage definitions and micro-batch flows
- Gradient accumulation and synchronization points
- Resource allocation and load balancing specifications
- Communication pattern optimization hints

Syntax Design Principles

The proposed syntax balances declarative clarity with imperative control:

The Tech National. *Embracing the Future: Infrastructure as Code with Standard Programming Languages*.

2024. <https://thetechnational.com/embracing-the-future-infrastructure-as-code-with->

TensorFlow Team. *Introduction to graphs and tf.function*. 2024. https://www.tensorflow.org/guide/intro_to_graphs

Data Science Central. *Understanding Dataflow graphs in TensorFlow*. 2019. <https://www.datasciencecentral.com/understanding-dataflow-graphs-in-tensorflow/>

ONNX Community. *ONNX | Home*. 2024. <https://onnx.ai/>

Encord Team. *Understanding ONNX: Enhancing AI Model Interoperability Across Platforms*. 2024. <https://encord.com/blog/onnx-open-neural-network-exchange-format/>

Multi-level abstraction enables progressive complexity disclosure while maintaining system-wide coherence.

```
# Declarative pipeline definition
pipeline "training_pipeline":
  stages: 4
  micro_batch_size: 32
  gradient_accumulation: true

  stage "embedding":
    resources: {gpu_memory: "8GB", compute: "V100"}
    layers: ["embedding", "attention_1"]

# Imperative optimization hints
optimize:
  - balance_computation_load()
  - minimize_communication()
```

Implementation Strategy

Standardization Roadmap

The implementation follows a three-phase approach:

Phase 1: Core Language Definition establishes syntax, semantics, and basic tooling infrastructure.

Phase 2: Pipeline Specialization adds native pipeline constructs and optimization algorithms.

Phase 3: Ecosystem Integration develops platform connectors and community governance.

Phased implementation enables iterative refinement while building ecosystem support progressively.

Validation Framework

Comprehensive validation encompasses conformance testing, semantic checking, and cross-platform compatibility verification.

Research Gaps and Future Directions

Identified Opportunities

Three primary opportunities emerge from the analysis:

Unified Infrastructure Abstraction addresses current tool fragmentation by providing end-to-end infrastructure modeling capabilities.

Performance-Aware Scheduling integrates performance models directly into infrastructure definitions for automatic optimization.

Multi-Cloud Portability enables vendor-neutral representation for true infrastructure portability.

Research Challenges

Key challenges include complexity management, performance modeling accuracy, and standardization process coordination across competing industry interests.

Conclusions

The research reveals a significant opportunity for developing a specialized graph language incorporating lessons from pipelined model parallelism. Current standards inadequately address temporal, resource-aware, and optimization needs of modern infrastructure.

The proposed infrastructure graph language would provide unified, performance-aware, and pipeline-native representation suitable for distributed systems management. This approach could significantly improve infrastructure efficiency, reliability, and developer productivity while establishing foundations for next-generation DevOps and MLOps practices.

As discussed in Section 2, the sophisticated scheduling and optimization techniques developed for ML pipeline training provide valuable patterns for general infrastructure orchestration. The implementation strategy outlined in Section 6 provides a roadmap for realizing these opportunities.

Future work should focus on prototype development, industry engagement, and standardization body coordination to realize these opportunities.

The integration of pipeline parallelism insights with infrastructure modeling represents a novel contribution to the field.

References

- [1] Microsoft DeepSpeed Team. Pipeline parallelism - DeepSpeed, 2024. URL <https://www.deepspeed.ai/tutorials/pipeline/>. DeepSpeed Documentation.
- [2] Deepak Narayanan, Aaron Harlap, Amar Phanishayee, Vivek Seshadri, Nikhil R. Devanur, Gregory R. Ganger, Phillip B. Gibbons, and Matei Zaharia. PipeDream: Generalized pipeline parallelism for DNN training. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles*, pages 1–15. ACM, 2019.

- [3] Insu Jang. Parallelism in distributed deep learning, June 2022. URL <https://insujang.github.io/2022-06-11/parallelism-in-distributed-deep-learning/>. Personal Blog.
- [4] Yanping Huang, Youlong Cheng, Ankur Bapna, Orhan Firat, Dehao Chen, Mia Xu Chen, HyoukJoong Lee, Jiquan Ngiam, Quoc V. Le, Yonghui Wu, and Zhifeng Chen. GPipe: Efficient training of giant neural networks using pipeline parallelism. In *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019.
- [5] Microsoft Research. PipeDream: A more effective way to train deep neural networks using pipeline parallelism, October 2019. URL <https://www.microsoft.com/en-us/research/blog/pipedream-a-more-effective-way-to-train-deep-neural-networks-using-pipeline-parallelism/>. Microsoft Research Blog.
- [6] Computer Science Stack Exchange. Scheduling distributed computational graph, 2015. URL <https://cs.stackexchange.com/questions/50243/scheduling-distributed-computational-graph>. Stack Exchange Discussion.
- [7] Wikipedia Contributors. DOT (graph description language), 2024. URL [https://en.wikipedia.org/wiki/DOT_\(graph_description_language\)](https://en.wikipedia.org/wiki/DOT_(graph_description_language)). Wikipedia.
- [8] Graphviz Team. DOT language, 2024. URL <https://graphviz.org/doc/info/lang.html>. Graphviz Documentation.
- [9] Wikipedia Contributors. GraphML, 2024. URL <https://en.wikipedia.org/wiki/GraphML>. Wikipedia.
- [10] ISO/IEC. ISO/IEC 39075:2024 - information technology — database languages — GQL. Iso standard, International Organization for Standardization, 2024. URL <https://www.iso.org/standard/76120.html>.
- [11] Joab Jackson. GQL: A new ISO standard for querying graph databases, June 2024. URL <https://thenewstack.io/gql-a-new-iso-standard-for-querying-graph-databases/>. The New Stack.
- [12] Neo4j Inc. Graph DB for network database infrastructure monitoring, 2024. URL <https://neo4j.com/use-cases/network-and-it-operations/>. Neo4j Use Cases.

- [13] Kubernetes Contributors. Cluster architecture, 2024. URL <https://kubernetes.io/docs/concepts/architecture/>. Kubernetes Documentation.
- [14] The New Stack. Reframing Kubernetes observability with a graph, July 2023. URL <https://thenewstack.io/reframing-kubernetes-observability-with-a-graph/>. The New Stack.
- [15] Apache Airflow. Dags — Airflow documentation, 2024. URL <https://airflow.apache.org/docs/apache-airflow/stable/core-concepts/dags.html>. Apache Airflow Documentation.
- [16] Apache Airflow. What is Airflow? — Airflow documentation, 2024. URL <https://airflow.apache.org/docs/apache-airflow/stable/index.html>. Apache Airflow Documentation.
- [17] The Tech National. Embracing the future: Infrastructure as code with standard programming languages, August 2024. URL <https://thetechnational.com/embracing-the-future-infrastructure-as-code-with-standard-programming-languages/>. The Tech National.
- [18] TensorFlow Team. Introduction to graphs and tf.function, 2024. URL https://www.tensorflow.org/guide/intro_to_graphs. TensorFlow Documentation.
- [19] Data Science Central. Understanding dataflow graphs in TensorFlow, June 2019. URL <https://www.datasciencecentral.com/understanding-dataflow-graphs-in-tensorflow/>. Data Science Central.
- [20] ONNX Community. ONNX | home, 2024. URL <https://onnx.ai/>. ONNX Official Website.
- [21] Encord Team. Understanding ONNX: Enhancing AI model interoperability across platforms, August 2024. URL <https://encord.com/blog/onnx-open-neural-network-exchange-format/>. Encord Blog.