Action Plan: Modeling the Dædælus Bandwidth and Reliability Argument

To: Dean

CC: Paul, Steve, Dugan, Kevin, Varun

Dean, (CC'd team),

Following up on your excellent progress with the Live Computation Model based on Metcalfe's work, it's time to formalize the narrative we want to build. The goal is to create a series of compelling, interactive models that visually and computationally prove our core thesis: the industry's reliance on bandwidth as the primary performance metric is fundamentally flawed. We need to show that focusing on the reliability of round-trip interactions is not only more efficient but also essential for the correctness of modern distributed systems.

This document outlines a multi-step plan to build this argument, starting with modeling historical concepts and culminating in a demonstration of the Dædælus architecture's advantages.

The Core Argument: Multiplexing Transactions, Not Bandwidth

Our central argument, as laid out in "Bandwidth Works in Practice, not in Theory," is a challenge to a 50-year-old assumption. As Metcalfe noted in 1976, when multiple stations share a medium, the bandwidth is divided among them. The industry has extended this reasoning to modern switched networks, leading to the belief that contention must be managed by fairly dividing bandwidth.

We want to prove this assumption is wrong for modern systems. The resource we should be multiplexing is not raw bandwidth, but the **transactional capacity of the link**. The fundamental unit of communication is the reliable, acknowledged, round-trip interaction. Our models should demonstrate this from first principles.

Modeling and Visualization Plan (Primarily in Mathematica)

Here are the four key models we need to build to tell this story.

1. Agent-Based Simulation of the Classical Half-Duplex Æthernet

First, we need to create a visual, agent-based simulation of the shared "Æther" as described in Metcalfe's 1976 paper. This must be a *computational* model, not merely statistical.

- **What to Model:**

    - Visualize multiple stations ("nodes") attempting to send "packets" over a single shared medium.
    - Implement the **Transmission and Contention Intervals** (Section 6 of Metcalfe's paper).
    - Computationally model the core mechanisms:
        - **Carrier and Interference Detection** (Sec 3.5).
        - **Collision Consensus Enforcement** (Sec 3.5).
        - **Binary Exponential Backoff** for retransmission (Sec 4.4).
- **Objective:** Visually demonstrate that as more stations contend for the medium, TCP throughput is decimated due to increasing round-trip latency. We want to create a timeline of the medium that clearly shows nodes colliding, backing off, and waiting, proving that contention fundamentally slows down the interaction rate.

2. The Problem on Modern Full-Duplex Links

Next, adapt the model to show how the same problem persists even on modern full-duplex links when bandwidth is the multiplexed resource.

- **What to Model:**
    - Create a "Bandwidth-Multiplexed Channel" where multiple, independent node pairs (A ↔ A, B ↔ B, etc.) communicate over a single, lossy (e.g., 0.1% packet loss) link.
- **Objective:** Show that even with separate transmit and receive paths, multiplexing the link's bandwidth still forces interactions into a chaotic, serialized queue. The result should be a similar degradation in round-trip latency and overall throughput, connecting directly to the Mathis equation (BW∝1/(RTT Sqrt(p))) and the graphs from Andy Helland's presentation. This proves that simply making the pipe full-duplex doesn't solve the core contention problem.

3. Demolishing the Stop-and-Wait Assumption with Point-to-Point Links

This is where we introduce the "circulating snakes" concept and show how our architecture escapes the limitations of classical models.

- **What to Model:**
    - A reliable, point-to-point link between exactly two nodes with a very short propagation delay (e.g., a few nanoseconds, as with a short Thunderbolt cable).
    - Model a continuous stream of bits (a "snake") that is longer than the physical link itself.
    - Provide a slider to adjust the link length.

- **Objective:** Prove that when a packet is "longer than the wire," the receiver can process the head of the packet and send an acknowledgment *while the sender is still transmitting the tail*. This demonstrates that we do not lose half the bandwidth to acknowledgments. The model should show that interaction latency approaches zero and link utilization approaches 100%, breaking the assumption that reliable stop-and-wait interactions must throttle throughput.

4. The Solution: Multiplexing Reliable Interactions

This final model will showcase the Dædælus approach as the superior solution.

- **What to Model:**
  - Use the same multi-pair setup as in Model #2 (A ↔ A, B ↔ B, etc.).
  - Instead of letting the flows fight for bandwidth, the channel should service them in a deterministic, round-robin fashion, interleaving one atomic handshake interaction from each pair.
- **Objective:** Provide definitive visual proof that "Interaction-Multiplexing" is fundamentally more efficient. The model should show that total throughput is maximized and that latency is deterministic because the protocol handles errors without costly timeouts and retransmissions. This is the function of our **Graph Virtual Machine (GVM)**—to deterministically schedule these interactions.

Immediate Next Steps & Schedule

- **This Week:** Focus on building out these models in Mathematica. Start with Model #1 to deeply understand the dynamics of the classical system.
- **Next Week:** Transition to coding. The goal is to implement **accurate slice-level emulation** in the DDL_Emulator. This means getting the "two snakes" state machine working, initially with 64-byte transfers and then moving to 8-byte transfers.
- **Reading:** Continue to critically analyze the **"Bandwidth Works in Practice, not in Theory"** paper and the foundational Metcalfe paper. Your insights from the models should directly inform our papers and specifications. We have an opportunity to submit our work to the journal *Entropy*, and these models will form the core of our argument.

These models are critical for our demos, our papers, and for telling the Dædælus story. Let's start by attacking Model #1.

Thanks,

Sahas